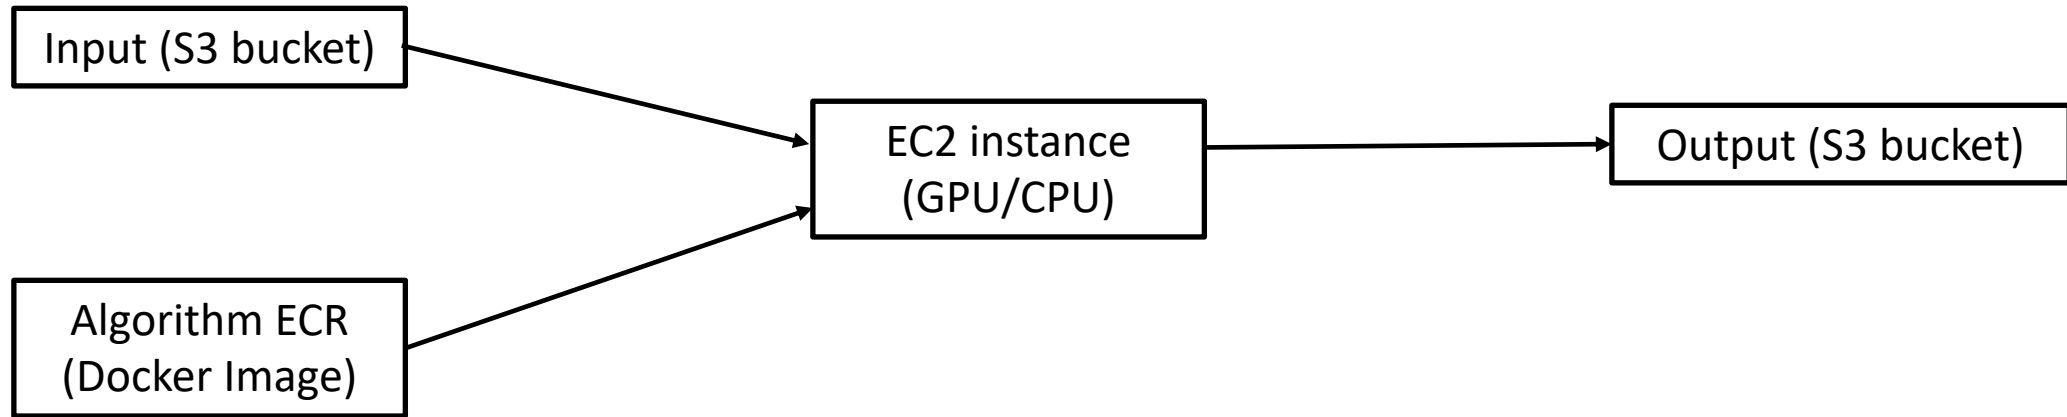


SageMaker Training Job in AWS

Abinash Pun

AISQUAD

Outline



- Amazon **S3** or **Amazon Simple Storage Service** is a service offered by Amazon Web Services that provides object storage through a web service interface.
- Amazon **Elastic Container Registry (Amazon ECR)** is an AWS managed container image registry service that is secure, scalable, and reliable.
- Amazon **Elastic Compute Cloud (EC2)** is the Amazon Web Service you use to create and run virtual machines in the cloud (we call these virtual machines 'instances').

Docker

- Software level virtualization
- creating image containing all the software dependencies to run the code
- **Build Docker**
 - Dockerfile: all the configuration
 - `$ docker build -t image_name .`
- **Push to ECR**
- ***build-and-push.sh***: script to build and push docker image
- ***ecr_image_fullname.txt***: contains image name with path (Do we always need to do this?)

Input and Output

- Initial Input: s3://../input
- E2:
 - Input: /opt/ml/input
 - Copy to /opt/ inside container (?)
 - Output: copy to /opt/ml/output
- Final output: s3://../output

Manual Setup: Amazon SageMaker > Training jobs > Create training job

- Job Setting
 - Job name:
 - IAM role: *SageMakerExecutionRole*
 - Algorithm Option
 - Algorithm source: *Your own algorithm container in ECR (?)*
 - Choose an Algorithm: *choose the image file (?)*
- Resource configuration
 - Instance type, Instance count, Additional storage volume per stance
 - Keep alive period (*waiting period?*)
 - Encryption key (for data)
- Stopping condition
 - Maximum runtime
- Network
 - Enable network isolation
 - VPC (Virtual Private Cloud)
- Input Data Configuration
 - Channels: train
 - Data source (S3), S3 Data type, S3 Data location
- Checkpoint configuration
- Output Data Configuration
- Managed Spot Training
- Tags

Setup via Script: *training_job.py*

```
1 import os
2 from urllib.parse import urlparse
3 import io
4 import boto3
5 import json
6 import jsonlines
7 import sagemaker
8 import pandas as pd
9 import numpy as np
10 from PIL import Image
11 from itertools import cycle, islice
12 import matplotlib.pyplot as plt
13 import matplotlib.patches as patches
14 from sagemaker.estimator import Framework, Estimator
15 from sagemaker.processing import Processor, ProcessingInput, ProcessingOutput
16
17 sagemaker_session = sagemaker.Session()
18
19 # we are using the notebook instance role for training in this example
20 #role = 'AmazonSageMaker-ExecutionRole-833537904510'
21 role = sagemaker.get_execution_role(sagemaker_session=sagemaker_session)
22 # you can specify a bucket name here, we're using the default bucket of SageMaker
23 bucket = "sagemaker-studio-833537904510-3lvvbxobayc"
24
25 print(role)
26 print(bucket)
27
28 with open(os.path.join('ecr_image_fullname.txt'), 'r') as f:
29     container = f.readlines()[0][:-1]
30
31 print(container)
32 s3_input = 's3://{}/auto_synth'.format(bucket)
33 s3_autobox_input = '{}/auto_box'.format(s3_input)
34 s3_synth_data_input = '{}/synth_data'.format(s3_input)
35
36 print(s3_autobox_input)
37 print(s3_synth_data_input)
38
39 cfg='{}/cfg/'.format(s3_input)
40 outpath='{}/results/'.format(s3_input)
41 images='{}/images/'.format(s3_autobox_input)
42 labels='{}/labels/'.format(s3_autobox_input)
43 background='{}/background/'.format(s3_synth_data_input)
44 bg_noise='{}/bg_noise/'.format(s3_synth_data_input)
```

```
45 bg_masks='{}/bg_masks/'.format(s3_synth_data_input)
46
47 print(cfg)
48 print(outpath)
49 print(images)
50 print(labels)
51 print(background)
52 print(bg_noise)
53 print(bg_masks)
54
55 # Job configuration
56 from sagemaker.session import TrainingInput
57
58 inputs = {
59     "cfg": TrainingInput(cfg),
60     "images": TrainingInput(images),
61     "labels": TrainingInput(labels),
62     "background": TrainingInput(background),
63     "bg_noise": TrainingInput(bg_noise),
64     "bg_masks": TrainingInput(bg_masks),
65 }
66
67 estimator = Estimator(
68     image_uri=container,
69     role=role,
70     instance_count=1,
71     instance_type='ml.m5.large',
72     # instance_type='local',
73     input_mode='File',
74     output_path=outpath,
75     base_job_name='auto-synth-test1'
76 )
77
78 estimator.fit(inputs)
```

Some Details

auto_box / Dockerfile [Info](#)

```
1 ARG BASE_IMG=${BASE_IMG}
2 #ARG BASE_IMG='pytorch/pytorch:latest'
3 FROM ${BASE_IMG}
4 RUN mkdir -p /opt
5 ENV PATH="/opt:${PATH}"
6 RUN apt-get update \
7     && apt-get install build-essential -y --no-install-recommends --allow-unauthenticated \
8     jq git gcc libgl1-mesa-dev wget gsutil libglib2.0-0 cmake
9 ## fix /usr/local/cuda-10.0/compat/libcuda.so
10 ## RUN bash -c 'echo "/usr/local/cuda-10.0/compat" > /etc/ld.so.conf.d/cuda.conf'
11 RUN ldconfig -v
12 #RUN pip install tensorboard torch torchvision --upgrade
13
14 COPY requirements.txt /opt/requirements.txt
15 COPY . /opt
16 RUN pip install -r /opt/requirements.txt
17 RUN cd /opt
18 WORKDIR /opt
19 RUN chmod +x train
20 ## https://github.com/aws/sagemaker-pytorch-training-toolkit/issues/143#issuecomment-566776288
21 ## https://github.com/aws/sagemaker-pytorch-training-toolkit/blob/upgrade-training-
22 toolkit/docker/build_artifacts/start_with_right_hostname.sh
23 ## https://github.com/aws/deep-learning-containers/blob/v2.0-pt-1.5.1-
24 py36/pytorch/training/docker/1.5.1/py3/Dockerfile.gpu#L181
25 COPY start_with_right_hostname.sh /usr/local/bin/start_with_right_hostname.sh
```

auto_box / start_with_right_hostname.sh [Info](#)

```
1 #!/usr/bin/env bash
2 if [[ "$1" = "train" ]]; then
3     CURRENT_HOST=$(jq .current_host /opt/ml/input/config/resourceconfig.json)
4     sed -ie "s/PLACEHOLDER_HOSTNAME/$CURRENT_HOST/g" changehostname.c
5     gcc -o changehostname.o -c -fPIC -Wall changehostname.c
6     gcc -o libchangephostname.so -shared -export-dynamic changehostname.o -ldl
7     LD_PRELOAD=/opt/code/libchangephostname.so train
8 else
9     eval "$@"
10 fi
```

Some Details

auto_box / build-and-push.sh [Info](#)

```
1 #!/bin/bash
2 image=$1
3 if [ "$image" == "" ]
4 then
5     echo "Usage: $0 <image-name>"
6     exit 1
7 fi
8 account=$(aws sts get-caller-identity --query Account --output text)
9 region=$(aws configure get region)
10 fullname="${account}.dkr.ecr.${region}.amazonaws.com/${image}:latest"
11 aws ecr describe-repositories --repository-names "${image}" > /dev/null 2>&1
12 if [ $? -ne 0 ]
13 then
14     aws ecr create-repository --repository-name "${image}" > /dev/null
15 fi
16 # Get the login command from ECR and execute it directly (for our own registry and for the base image registry)
17 aws ecr get-login-password --region ${region} | docker login --username AWS --password-stdin
   ${account}.dkr.ecr.${region}.amazonaws.com
18 base_img='pytorch/pytorch:latest'
19 echo 'base_img:'$base_img
20 # Build the docker image locally and then push it to ECR with the full name.
21 cd container
22 echo "Building image with name ${image}"
23 docker build --no-cache -t ${image} -f Dockerfile --build-arg BASE_IMG=$base_img .
24 docker tag ${image} ${fullname}
25 echo "Pushing image to ECR ${fullname}"
26 docker push ${fullname}
27 # Writing the image name to let the calling process extract it without manual intervention:
28 echo "${fullname}" > ecr_image_fullname.txt
```